

# Linux Kernel Training: Lecture 1

Building the Software for BEAGLEBONE BLACK

---

Sam Protsenko

March 21, 2020

GlobalLogic

# Agenda

1. Hardware Overview
2. Software Overview
3. Perspective on Building

# Organization

# Linux Kernel ProCamp Details

- Mon/Fri, 3pm - 5pm
- Schedule: [link](#)
- Target: BEAGLEBONE BLACK and QEMU
- Host: Personal laptop (Ubuntu 18.04) or Training Centre PC
- Training Centre PC:
  - Press F9 on boot (show boot menu)
  - Select second drive (TS64GSSD370S, 64 GB)
  - Login: Lin-Ker
  - Password: 123

- Oleksandr Redchuk <[oleksandr.redchuk@globallogic.com](mailto:oleksandr.redchuk@globallogic.com)>
- Sam Protsenko <[semen.protsenko@globallogic.com](mailto:semen.protsenko@globallogic.com)>
- Andrii Lukin <[andrii.lukin@globallogic.com](mailto:andrii.lukin@globallogic.com)>
- Vitaliy Vasylyskyy <[vitaliy.vasylyskyy@globallogic.com](mailto:vitaliy.vasylyskyy@globallogic.com)>
- Illia Smyrnov <[illia.smyrnov@globallogic.com](mailto:illia.smyrnov@globallogic.com)>
- Ruslan Bilovol <[ruslan.bilovol@globallogic.com](mailto:ruslan.bilovol@globallogic.com)>

# Hardware Overview

---

# Embedded Systems

## Embedded Systems:

- Mobile
- Automotive
- Networking
- Smart TVs, game consoles, set-top boxes
- IoT
- Medical
- Aerospace
- Industry



# Embedded Programming

Differences from regular system:

- Cross-compiling
- Flashing
- Serial console
- Testing concerns
- Working with hardware
- Non discoverable buses on board (device tree, platform drivers)





# ARM Cortex A/R/M Families


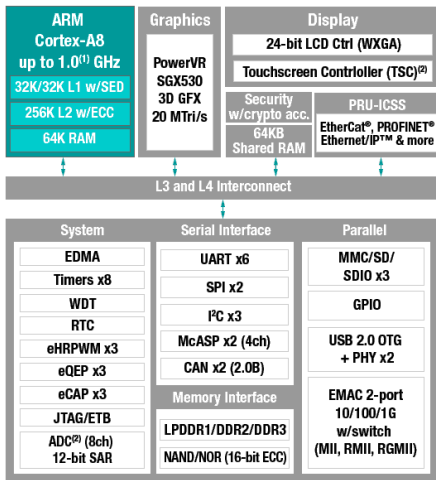


Figure 1: ARM Cortex Processor Families


# AM335x SoC



- **Board:** BEAGLEBONE BLACK (Texas Instruments)
- **SoC:** AM3358 (Texas Instruments)
- **Processor core:** Cortex-A8
- **Processor architecture:** ARMv7-A

Figure 2: AM335x Functional Diagram

# BeagleBone Black



# BeagleBone Black: Pros and Cons

## Pros:

- Open Hardware
  - Public TRM
  - Schematic
  - PCB files
- Supported in upstream
  - Kernel
  - U-Boot
- Conventional ARM architecture
- Very popular
- Low cost (\$55)

## Cons:

- Old 32-bit architecture
- Single core processor
- Android is not supported officially
- No WiFi

# Software Overview


---

# Software Components

- U-Boot
- Linux kernel
- BusyBox


# Software Components

- U-Boot
- Linux kernel
- BusyBox



# Software Components


- U-Boot
- Linux kernel
- BusyBox





# Software Components

- U-Boot
- Linux kernel
- BusyBox




## Building Steps

1. Obtain the software
2. Checkout to desired branch or tag
3. Consult with **README** and **INSTALL**
4. Install all build dependencies
5. Configure shell environment for cross-compiling
6. **Configure the software for build with options you desire**
7. **Build the software**
8. **Install/flash the built software**

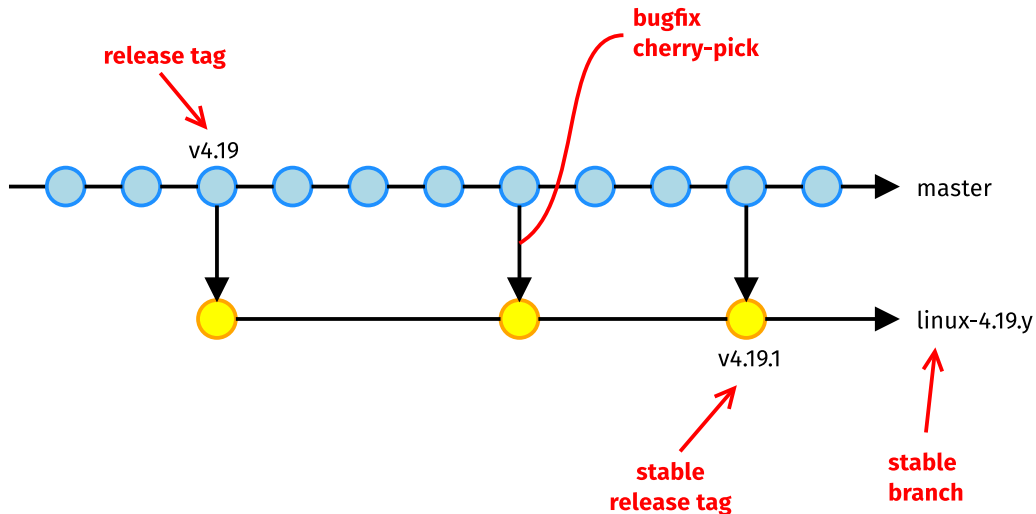
## Perspective on Building

---

# Kernel Branching Strategy




# Kernel Branching Strategy



- Git tags: for releases (e.g. `v4.19`)
- Git branches: for stable releases (e.g. `linux-4.19.y`)
- Some stable branches are LTS (Long Term Support)
- When possible, let's use stable branches (for reliability)
- When stable branches are not available, let's use release tags

## Toolchain (page 1)



- Set of tools for cross-compiling:
  1. `gcc`
  2. binutils: `ld`, `as`, `objdump`, `objcopy`, `readelf`, etc.
  3. `glibc` and other system libraries (optional)
  4. Linux kernel headers (optional)
  5. `gdb` (optional)



- Set of tools for cross-compiling:
  1. `gcc`
  2. binutils: `ld`, `as`, `objdump`, `objcopy`, `readelf`, etc.
  3. `glibc` and other system libraries (optional)
  4. Linux kernel headers (optional)
  5. `gdb` (optional)
- Toolchain types:
  - Bare-metal targeted (`arm-eabi`): for U-Boot and kernel
  - Linux targeted (`arm-linux-gnueabihf`): for BusyBox
- In our case: host = x86\_64, target = ARMv7-A
- We use only GNU (GCC-based) toolchains

Toolchain tuple examples:

- **arm-foo-none-eabi**, bare-metal toolchain targeting the ARM architecture, from vendor *foo*
- **arm-unknown-linux-gnueabi**, Linux toolchain targeting the ARM architecture, using the EABI ABI and the glibc C library, from an *unknown* vendor
- **armeb-linux-uclibcgnueabi**, Linux toolchain targeting the ARM big-endian architecture, using the EABI ABI and the uClibc C library
- **mips-img-linux-gnu**, Linux toolchain targeting the MIPS architecture, using the glibc C library, provided by Imagination Technologies

- Regular compilation on host system (using *native toolchain*):

```
$ gcc main.c
```

- Regular compilation on host system (using *native toolchain*):

```
$ gcc main.c
```

- Cross-compilation for ARM target (part before `gcc` is called *prefix*):

```
$ /toolchain/path/bin/arm-eabi-gcc main.c
```

- Regular compilation on host system (using *native toolchain*):

```
$ gcc main.c
```

- Cross-compilation for ARM target (part before `gcc` is called *prefix*):

```
$ /toolchain/path/bin/arm-eabi-gcc main.c
```

- More universal way:

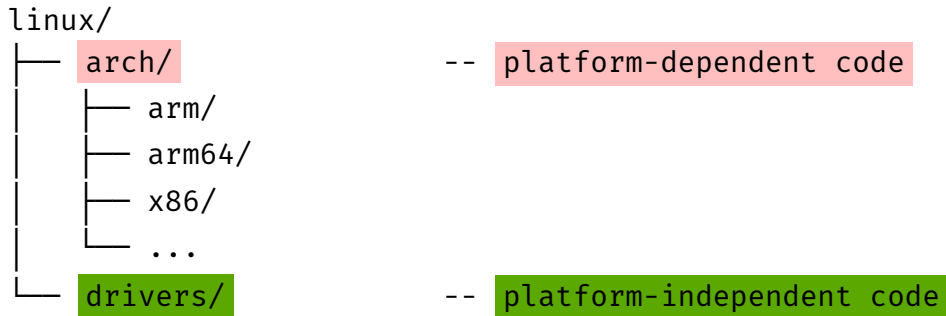
```
$ PATH=/toolchain/path/bin:$PATH
```

```
$ CROSS_COMPILE=arm-eabi-
```

```
$ ${CROSS_COMPILE}gcc main.c
```

# Specifying Architecture

- Kernel supports many CPU architectures:



- We need to choose which architecture to build for:

```
$ export ARCH=arm
```

- Shell environment configuration for building U-Boot/kernel/BusyBox:

```
$ export ARCH=arm
```

```
$ export PATH=/toolchain/path/bin:$PATH
```

```
$ export CROSS_COMPILE=arm-eabi-
```

- Makefile utilizes those env vars

Take Five



## Kbuild: User's Perspective

## Building: General Steps

- All projects (U-Boot, Linux kernel and BusyBox) use Kbuild
- Build steps: configuration, build, installation
- **Configuration** (generate `.config` file):

```
$ make defconfig
```

- **Build:**

```
$ make
```

- **Installation:**

```
$ make install
```

## Building: Custom Configuration

- Sometimes existing `defconfig` is not enough
- How can we customize our configuration?
  - Using `make menuconfig`
  - Using `merge_config.sh` script
  - Using old `.config` file

## Building: Custom Configuration

- Sometimes existing `defconfig` is not enough
- How can we customize our configuration?
  - Using `make menuconfig`
  - Using `merge_config.sh` script
  - Using old `.config` file
- Example: kernel configuration using `merge_config.sh`:

```
$ ./scripts/kconfig/merge_config.sh \
  arch/arm/configs/multi_v7_defconfig \
  fragments/bbb.cfg
```

## Building: .config Example

Excerpt from `.config` file:

```
CONFIG_USE_OF=y
CONFIG_DEFAULT_HOSTNAME="(none)"
CONFIG_CMDLINE=""
# CONFIG_PREEMPT is not set
CONFIG_I2C_GPIO=m
CONFIG_LOG_BUF_SHIFT=17
# CONFIG_SLAB is not set
CONFIG_USB=y
CONFIG_SND_USB_AUDIO=m
```


## Kernel Modules

- Every driver is a module
- Kernel modules can be:
  - Loadable: “=m”
  - Built-in: “=y”
- Kernel loadable module (**.ko** file) is some code that can be loaded into kernel space (i.e. added to running kernel as a plugin)

# Kernel Modules

- Every driver is a module
- Kernel modules can be:
  - Loadable: “=m”
  - Built-in: “=y”
- Kernel loadable module (**.ko** file) is some code that can be loaded into kernel space (i.e. added to running kernel as a plugin)
- How it works:
  - **multi\_v7\_defconfig** is common for all ARMv7 systems (so the single **zImage** can be used)
  - Device Tree file covers SoC and board differences
  - Needed modules (for particular board) can be loaded in run-time
- It's not always convenient to load a lot of modules

## How to Speed-Up the Build? (page 1)





## How to Speed-Up the Build? (page 2)

- **Kbuild tracks all dependencies very well!**

- *Clean build* (try to avoid it):

```
$ make  
// Do some changes to source code  
$ make distclean  
$ make
```

- Use *incremental build* instead:

```
$ make  
// Do some changes to source code  
$ make
```

## How to Speed-Up the Build? (page 3)

- Distribute the compilation between CPU cores using multi-threading build:


```
$ make -j4
```

- Use **ccache** tool:


- Caches **.o** files on the first build
- On next build, if some **.o** files are unchanged, cached versions will be used
- **ccache** creates a hash by **.c** file content and by build command
- ...So if you change the toolchain, cache won't be used
- Speed up for clean build is usually 5 times
- Can be used as a wrapper:

```
$ ccache gcc main.c
```

## How to Speed-Up the Build? (page 4)



## How to Speed-Up the Build? (page 5)



RootFS

# RootFS

What is RootFS?

- Filesystem that is needed to make userspace work
- Mounted to “/”
- Crucial component is `init` tool
- Besides of that: libc, kernel modules, tools, config files...

# RootFS


What is RootFS?

- Filesystem that is needed to make userspace work
- Mounted to “/”
- Crucial component is `init` tool
- Besides of that: libc, kernel modules, tools, config files...

Known rootfs's for BBB:


- Debian
- Yocto/OpenEmbedded
- BuildRoot
- BusyBox

## BusyBox Linking: Static





## BusyBox Linking: Dynamic



## BusyBox Linking: Static vs Dynamic

- **Static linking:** libc (.a) is compiled in your binary
  - Only “**busybox**” binary is needed in rootfs
  - Easier to build and minimal
  - Some networking functions won't work (like **nslookup**, see libnss)
- **Dynamic linking:** libc used as a shared library (.so)
  - Only one copy of libc is used (for all possible apps)
  - Dynamic libraries must be copied in rootfs **/lib** (libc and its dependencies)

## BusyBox Applets

- BusyBox is a multi-call binary
- Apps in BusyBox rootfs are just symbolic links:

/bin

```
|— busybox  
|— grep -> busybox  
|— ls    -> busybox
```

- So you can call **ls** tool like this:

```
# busybox ls -l
```

- ...which is identical to this form:

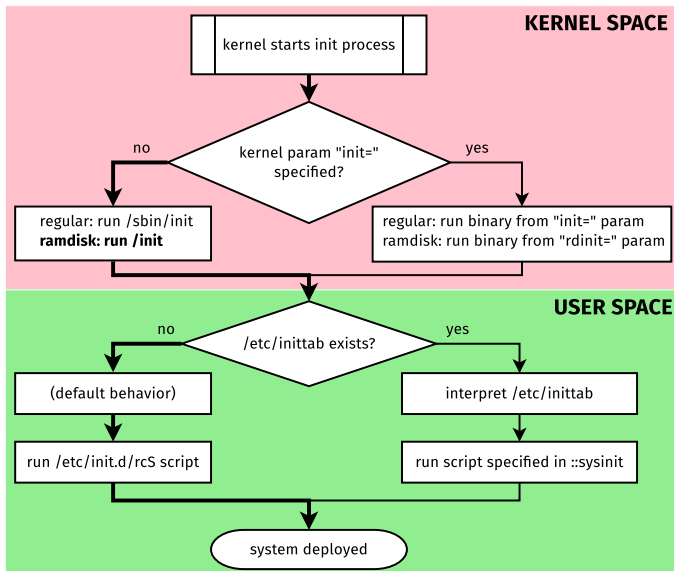
```
# ls -l
```

# Init process

- First process started during boot (kernel starts it)
- PID = 1, uid = 0 (root)
- It never exits (daemon)
- Init is a parent for all processes
- Automatically adopts all orphaned processes
- Init is started by the kernel using a hard-coded filename (e.g. `/init`)
- A kernel panic will occur if the kernel is unable to start it
- Most popular init implementations:
  - sysvinit
  - openrc
  - upstart
  - systemd

- **busybox** tool implements init as an applet
- BusyBox's init implementation resembles SysVinit, but more simple
- Doesn't support runlevels (as opposed to SysVinit)
- Starts **/etc/init.d/rcS** script
- (Re)spawns children according to **/etc/inittab** (e.g. **getty**)
- Handles signals (e.g. **reboot** and **poweroff**)

## BusyBox init (cont'd)



- No **udev** in BusyBox
- *'mdev is a mini-udev implementation for dynamically creating device nodes in the /dev directory'*
- Requires SysFS support in kernel; it must be mounted to **/sys**
- Can be also used for hot-plugging (e.g. load needed kernel module when some USB device was inserted)
- **mdev -s**: scan **/sys** and populate **/dev**
- **mdev** without params: kernel hotplug helper
- For more details see: **doc/mdev.txt**

## BusyBox init: Script Example

- **rc** = “run commands”, **S** = single-user runlevel
  - Example of `/etc/init.d/rcS` file:
- 

```
#!/bin/sh
```

```
mount -t sysfs none /sys
```

```
mount -t proc none /proc
```

```
mount -t debugfs none /sys/kernel/debug
```

```
echo /sbin/mdev > /proc/sys/kernel/hotplug
```

```
mdev -s
```



## BusyBox rootfs (static, minimal)

```
/
├── bin
│   └── busybox
├── boot
│   ├── am335x-boneblack.dtb
│   └── zImage
├── dev
├── etc
│   ├── init.d
│   └── rcS
```

```
├── init -> bin/busybox
├── linuxrc -> bin/busybox
├── proc
├── root
├── sbin
├── sys
├── tmp
└── usr
    ├── bin
    └── sbin
```

Demo: menuconfig

# Assignments

---

# Assignment

- Using BBB instructions guide (will be sent out):
  - Go through 1st chapter (“Preparing the Tools”)
  - Go through 2nd chapter (“Obtaining and Building the Software”)
  - Run built software on QEMU, using section 3.1 “QEMU Boot”
- Download TRM and datasheet for AM335x
- Download schematic for BBB
- Proof: send me screenshot of `uname -a` output in your QEMU





## Advanced assignment (optional)

- Using TRM, figure out:
  - Which module (TRM section?) is used for setting clocks (gating, DPLL)
  - Which module (TRM section?) is used for pin multiplexing
  - Where GPIO output registers are documented
  - Where UART RX/TX registers are documented
- Using schematic, figure out:
  - Which pins (pads) the user LEDs are connected to
  - How to mux those pins for GPIO (use datasheet and TRM)?
  - Which registers to use for pin muxing and then blinking some LED?

## References

---

## Recommended Reading

-  Karim Yaghmour, Jon Masters and others.  
*Building Embedded Linux Systems.*
-  Brian Ward.  
*How Linux Works*, 2nd Edition.
-  Andrew N. Sloss and others.  
*ARM System Developer's Guide.*
-  Robert Love.  
*Linux Kernel Development*, 3rd Edition.

Thank you!